Governance Score Card

What your governance scan means.
One page. Five properties. One next step per gap.

EVOKED.DEV

Version 1.0 · 2026

Governance Score Card

Governance Score Card

How to Use This Card

Run `npx agent-governance-check` in your project directory. Match each result to the grid below.

Score Interpretation Grid

1. Decision Memory

FOUND - Your agent records its decisions. It can explain why it chose what it chose. This is the foundation of accountability.

PARTIAL - Some form of state persistence exists, but it may not be structured for governance. Check whether decisions are logged with reasoning, not just cached for performance.

MISSING - Your agent has no memory of its own decisions. Every invocation is a blank slate. Risk: no accountability trail. No way to audit what happened or why. If a regulator asks "why did your agent do this?" you cannot answer.

Next step: Add a `memory-schema.md` to your project. Define what the agent remembers, what it forgets, and who decides. Run `npx agent-governance-check --init` to get a starter template.

---

2. Refusal Capability

FOUND - Your agent has defined boundaries - things it will not do regardless of instruction. This is the hardest governance property to implement and the most important.

PARTIAL - Some guardrails exist (input filtering, content safety) but they may not be agent-level refusal. Guardrails filter inputs. Refusal is the agent's own judgment about what it should not do.

MISSING - Your agent will attempt anything it's asked. Risk: no ethical boundaries. No principled "no." If a user or upstream system sends a harmful instruction, your agent has no basis for refusing. The EU AI Act (Article 14) requires human oversight mechanisms - refusal is how agents participate in that.

Next step: Create a `restraint-spec.md` defining what your agent refuses and why. Six categories to consider: scope (outside expertise), safety (potential harm), consent (unauthorized actions), integrity (dishonest outputs), energy (resource abuse), and legal (regulatory violations).

---

3. Persistent Identity

FOUND - Your agent has a defined identity - a charter, persona, or role definition that persists across sessions. It knows what it is and what it's for.

PARTIAL - Some identity exists (system prompts, role descriptions) but it may not be structured as a governance artifact. A system prompt that changes with every deploy is configuration, not identity.

MISSING - Your agent has no defined identity beyond its prompt. Risk: identity drift. The agent's behavior changes silently as prompts evolve, models update, or context shifts. No way to measure whether the agent is still the agent it was designed to be.

Next step: Write an `agent-charter.md` that defines: who this agent is, what it's for, what its boundaries are, and what it values. This is the document you test against when measuring drift.

---

4. Disagreement Protocol

FOUND - Your system has a defined process for when agents disagree, when a human overrides an agent, or when conflicting directives arise. This is governance for the edge cases.

PARTIAL - Some conflict handling exists (error fallbacks, retry logic) but it may not address principled disagreement. An agent that silently defers to the last instruction is not disagreeing -

it's complying.

MISSING - No process for disagreement exists. Risk: silent failures. When agents receive conflicting instructions or encounter ethical edge cases, they fail unpredictably. In multi-agent systems, this compounds. The EU AI Act (Article 72) requires governance mechanisms for AI systems - disagreement protocols are how you demonstrate them.

Next step: Create a `deliberation-framework.md` defining how conflicts are surfaced, who arbitrates, and how resolutions are recorded. Even single-agent systems need a process for "what happens when the human and the agent disagree."

---

5. Drift Detection

FOUND - You measure whether your agent's behavior has changed over time. Baselines exist. Monitoring is active. You would know if the agent drifted.

PARTIAL - Some monitoring exists (performance metrics, error rates) but it may not measure behavioral drift. Latency monitoring tells you the agent is fast. It doesn't tell you the agent is still aligned.

MISSING - No drift detection exists. Risk: silent degradation. Model updates, context changes, prompt evolution, and data shifts all cause behavioral drift. Without measurement, you won't know until a user reports a problem - or a regulator does.

Next step: Define behavioral baselines in a `drift-monitoring.md` file. What does "normal" look like for this agent? Measure against that baseline on a cadence (weekly, monthly). If any dimension exceeds your threshold, investigate.

What Your Score Means

Score	Assessment
5/5	Your agent has governance architecture. Maintain it. Measure it. You're ahead of 95% of the industry.

Score	Assessment
3-4/5	Foundation exists. The gaps are specific and addressable. Use the next-step actions above.
1-2/5	Governance is absent or minimal. Start with identity (charter) and refusal (restraint spec). Everything else builds on those two.
0/5	Your agent operates without governance. This is common in 2025-2026. It will not be acceptable by 2027. Start now.

Regulatory Context

The EU AI Act takes effect August 2, 2026. Key articles relevant to agent governance:

- **Article 14:** Human oversight - requires mechanisms for human intervention. Refusal capability is how agents participate.
- **Article 50:** Transparency obligations - requires disclosure of AI system capabilities and limitations. Identity and charter serve this.
- **Article 72:** Governance mechanisms - requires documented governance for AI systems. All five properties contribute.

Next Steps

If you want...	Go to...
Templates with filled-in examples	Guided Templates ($5)
Step-by-step implementation for your framework	Implementation Guide ($9.99)
The complete governance framework	Governance Starter Kit ($49)

If you want...	Go to...
Hands-on help with your system	Book a discovery call

Built by Evoked. evoked.dev

"We evoke - we never extract."