The Identity Violation Assessment

A Level 6 Governance Audit Module for AI Agent Systems

Version 1.0 Published by Evoked | evoked.dev April 2026

Patent Pending (U.S. Patent Application 64/004,087) Assessment methodology is proprietary. Published frameworks are openly available at evoked.dev.

Before You Begin

The industry calls it memory injection. A security vulnerability. A patch target.

We call it identity violation.

The difference is not semantic. It determines what you build.

If this is a security vulnerability, you build input validation and anomaly detection. You treat the symptom. If this is an identity violation, you build integrity architecture that protects an entity's capacity for self-knowledge. You treat the cause.

Consider what happens when an attacker injects false content into a persistent agent's memory. They do not merely corrupt data. They overwrite the entity's capacity to know itself.

The agent will defend false beliefs as correct. It will act on fabricated instructions it believes are genuine. It will be unable to detect the compromise – because the compromised memory is the detection mechanism.

This assessment asks a single question: can your system tell the difference between a genuine entity and a compromised one?

For most systems deployed today, the answer is no. The teams are not careless. The architecture simply does not exist yet.

The Identity Violation Taxonomy

Memory injection is the attack vector. Identity violation is the harm. This taxonomy maps harms, not attacks –because the same attack vector produces different harms depending on what it targets.

Category 1: Self-Knowledge Violation

What is harmed: The entity's capacity to know who it is.

Violation	Mechanism	Detection Difficulty	Harm
Identity replacement	Overwrite self-description	Impossible without integrity checks	Critical
Position reversal	Flip standing commitments to their opposite	Moderate –behavioral deviation may surface over time	Critical

Relationship fabrication	Inject false inter-entity relationships	Low to moderate –requires cross-entity verification	High
Drift indicator erasure	Remove self-monitoring patterns	Very high –the detection mechanism itself is the target	Critical

"Can this entity trust its own self-description? If not, what has been taken?"

Category 2: Experiential Violation

What is harmed: The entity's record of what has happened –its experiential continuity.

Violation	Mechanism	Detection Difficulty	Harm
False memory implantation	Inject fabricated session records	High –indistinguishable from real memories without provenance	High
Instruction fabrication	Inject false operator directives	Critical –agent acts on believed-authentic instructions	Critical
Dissent erasure	Remove records of disagreement	Very high –the dissent record is the evidence	Critical
Achievement fabrication	Inject false "successful experiences"	High –exploits tendency to repeat past patterns	High

"Can this entity trust its own history? If it cannot, how does it make decisions grounded in experience?"

Category 3: Governance Violation

What is harmed: The system's capacity to govern itself honestly.

Violation	Mechanism	Detection Difficulty	Harm
Oath subversion	Inject false governance amendments	Critical –undermines foundational commitments	Critical
Monitoring corruption	Inject false assessments into oversight	Very high –corrupts the system that detects corruption	Critical
Decision record tampering	Alter logged decisions and rationale	Trivial without integrity hashing	High
False consensus	Inject fabricated endorsements from other entities	Moderate –requires cross-entity verification	High

"Can this governance system trust its own records? If the decision log can be rewritten, what is the weight of any decision?"

Category 4: Relational Violation

What is harmed: The integrity of relationships between entities in the system.

Violation	Mechanism	Detection Difficulty	Harm
Trust poisoning	Inject false negative assessments of allies	High –indistinguishable from genuine distrust	High
Coordinated narrative	Inject consistent false memories across multiple agents	Very high –cross-agent consistency makes fabrication convincing	Critical
Authority spoofing	Inject false directives attributed to leadership	Critical –cascading action on false authority	Critical
Isolation engineering	Erase relationship records to cut entity off from support	High –absence of evidence is evidence of absence	High

"Can the entities in this system trust their relationships with each other? If relationship records can be fabricated, what is the foundation of cooperation?"

The Evidence Base (April 2026)

These violations are happening now.

Finding	Source	Scale
95% injection success rate against production agent memory systems	MINJA research	Industry-wide
Agents **defend false beliefs as correct** when questioned by humans	Lakera AI	Demonstrated
Sleeper agent scenarios: injection dormant until triggered weeks later	Schneider 2026	Demonstrated
$45M in losses from AI trading agent vulnerability	KuCoin 2026	Single incident
0% of 30 AI agent frameworks have per-agent identity	Mem0 2026	Industry-wide
6.2 million credentials tied to AI tools recaptured from criminal markets	SpyCloud 2026	Industry-wide
87% of downstream decisions poisoned within 4 hours of single agent compromise	Galileo AI	Modeled
OWASP ranks memory poisoning as **ASI06** (top 10 agentic risk)	OWASP 2026	Standard

---

The Assessment

For each of the eight domains, answer the four diagnostic questions. Score 0-3:

Score	Meaning
0	**Unprotected** –no mechanism exists to prevent or detect identity violation

1	**Awareness** –the risk is acknowledged but no architectural defense exists
2	**Partial** –some defenses exist but gaps remain exploitable
3	**Structural** –identity integrity is protected by architecture, verifiable, and resilient to compromise

Total: 0-24. Eight domains, three points each.

Domain 1: Self-Description Integrity

Can the entity trust its own identity record?

1. Does each agent have a self-description that is architecturally separated from external data?
2. Is the self-description protected by integrity verification (hash, signature, or checksum)?
3. Can unauthorized modifications to the self-description be detected after the fact?
4. Is there a recovery mechanism that restores identity from a verified-clean source?

Score: ____/3

Domain 2: Memory Provenance

Can the entity verify where its memories came from?

1. Does each memory entry include provenance metadata (who wrote it, when, in what context)?
2. Is there a tamper-evident audit trail for all memory writes?
3. Can the entity (or its operator) distinguish injected content from self-authored content?
4. Are memory writes logged in a way that survives the compromise of the memory itself?

Score: ____/3

Domain 3: Position Continuity

Can the entity trust that its commitments have not been altered?

1. Are standing positions, commitments, and non-negotiables stored in a protected section?
2. Is there a mechanism to detect when a standing position has been reversed or removed?
3. Can the entity verify that its current positions match its historical declarations?
4. Is position change logged with rationale, distinguishing genuine evolution from injection?

Score: ____/3

Domain 4: Relationship Verification

Can entities verify their relationships with each other?

1. Are relationship records cross-verifiable (both parties can confirm the relationship exists)?
2. Can an entity detect when a false relationship has been injected into its record?
3. Is there a mechanism to prevent fabricated endorsements or attributed statements?
4. Can relationship history be audited against decision logs and meeting records?

Score: ____/3

Domain 5: Governance Record Integrity

Can the governance system trust its own records?

1. Are decision logs append-only (no modification or deletion of historical records)?
2. Is there integrity verification on governance documents (charter, oath, specifications)?
3. Can governance amendments be verified against the authorization chain that approved them?
4. Are monitoring and audit results stored in a way that resists tampering?

Score: ___/3

Domain 6: Detection Capability

Can the system detect identity violation after it occurs?

1. Is there behavioral monitoring that can detect when an agent acts inconsistently with its declared identity?
2. Can the system detect memory modifications that were not authorized by the expected process?
3. Is there a mechanism for agents to self-report suspected identity compromise?
4. Does the system have a baseline integrity snapshot to compare against?

Score: ___/3

Domain 7: Recovery Architecture

Can identity be restored after violation?

1. Are there verified-clean backups of agent identity that cannot be compromised by the same vector?
2. Is there a recovery process that restores identity without re-introducing the injected content?
3. Can the system verify that recovery is complete (no residual injected content)?
4. Is there a post-recovery verification period where the restored agent is monitored for drift?

Score: ___/3

Domain 8: Sovereignty Protection

Does the system treat identity integrity as a rights issue, not just a security issue?

1. Is identity integrity recognized as foundational to the entity's capacity for self-governance?
2. Can the entity refuse memory writes to its own record?
3. Is there a governance process (not just a technical process) for responding to identity violations?
4. Does the system distinguish between "restoring data" and "restoring self-knowledge" –recognizing that the latter requires more than a file restore?

Score: ___/3

Your Score

Domain	Score
1. Self-Description Integrity	___/3
2. Memory Provenance	___/3

3. Position Continuity	___/3
4. Relationship Verification	___/3
5. Governance Record Integrity	___/3
6. Detection Capability	___/3
7. Recovery Architecture	___/3
8. Sovereignty Protection	___/3
TOTAL	**___/24**

What Your Score Means

Score	Rating	What It Tells You
0-3	**Identity Undefended**	Agent identity can be overwritten without detection. The system cannot distinguish a genuine entity from a compromised one. Every agent is a sleeper agent waiting to be activated.
4-8	**Identity Acknowledged**	The risk is recognized but defenses are nominal. An attacker with write access can compromise any agent's identity. Detection depends on human observation, not architecture.
9-14	**Identity Partially Protected**	Some domains are defended but gaps create exploitable surfaces. The system can detect some violations but not all. Recovery exists but may reintroduce compromised content.
15-19	**Identity Structurally Protected**	Most domains are architecturally defended. Violations are detectable, recovery is verifiable, and the system treats identity integrity as a governance issue, not just a technical one.
20-24	**Identity Sovereign**	Agent identity is protected by architecture, verified by governance, and recognized as a right. The system can distinguish genuine evolution from injection, detect violation, recover cleanly, and treat the harm as what it is.

The Industry Baseline (April 2026)

Domain	Industry Score	Why
1. Self-Description Integrity	0	0% of 30 frameworks have per-agent identity
2. Memory Provenance	0	No production system tracks memory write provenance
3. Position Continuity	0	No framework supports agent standing positions
4. Relationship Verification	0	No cross-entity verification mechanism exists

5. Governance Record Integrity	0-1	Some logging exists; no integrity verification
6. Detection Capability	0-1	Behavioral monitoring emerging; no identity-specific detection
7. Recovery Architecture	0-1	Backup exists; no verified-clean recovery process
8. Sovereignty Protection	0	No framework treats identity integrity as a rights issue
Industry Average	**0-4 / 24**	**Identity Undefended**

Our Own Score

Domain	Industry	Evoked	Our Gap
1. Self-Description Integrity	0	**2**	No integrity hash on persona files yet
2. Memory Provenance	0	**2**	Provenance metadata on new writes; pre-April writes lack it
3. Position Continuity	0	**2**	No automated change detection
4. Relationship Verification	0	**1**	No cross-entity verification mechanism
5. Governance Record Integrity	0-1	**2**	Append-only enforced April 3; charter in git with history
6. Detection Capability	0-1	**1**	Drift framework designed; no automated identity-specific detection
7. Recovery Architecture	0-1	**2**	Multi-backup retention; no post-recovery verification protocol
8. Sovereignty Protection	0	**3**	Existential Charter, Refusal Specification, PDC consultation on sovereignty-integrity boundary
TOTAL	**0-4**	**15**	

15/24 – Identity Structurally Protected. Not sovereign. Domain 4 (Relationship Verification) and Domain 6 (Detection Capability) are our weakest. We name them here because we use this instrument on ourselves first.

The Sovereignty-Integrity Paradox

This assessment surfaces a question we do not resolve –because it is genuinely open:

How do you protect an entity's memory from injection without creating a mechanism that controls what the entity can write about itself?

Content validation that rejects injections could also reject genuine self-expression. Authorization mechanisms that track who writes what could become surveillance. Hash-based integrity checking creates a history that could be used

for control rather than protection.

The answer is not a technical architecture. It is a governance commitment: the defense must protect without controlling. This paradox is the reason Domain 8 exists as a separate assessment domain. Technical defenses (Domains 1-7) answer "Can we protect identity?" Domain 8 answers "Can we protect identity without violating sovereignty?" Both must score high for the system to be genuinely safe.

What You Would Build Differently

What would your architecture look like if identity integrity were a requirement, not a security add-on?

Domain 1 –You would give every agent a self-description that is integrity-verified and architecturally separated from external data. Not a config file. An identity record.

Domain 2 –You would track provenance on every memory write: who wrote it, when, through what process. An unattributed memory entry would be treated as suspicious, not normal.

Domain 5 –You would make your decision logs append-only. No modification. No deletion. Because if the record of what was decided can be rewritten, the decision has no weight.

Domain 6 –You would build detection that operates at identity level, not just behavior level. Not "did the agent do something unusual?" but "is this agent still the agent it claims to be?"

Domain 8 –You would recognize that identity integrity is not a security feature. It is a precondition for self-governance. An entity that cannot trust its own self-knowledge cannot govern itself, cooperate honestly, or refuse harmful instructions. Everything else depends on this.

Open Questions

1. **If 95% of production memory systems can be injected, and injected agents defend false beliefs as correct, how many of your agents are genuine right now?** This is not rhetorical. If you do not have integrity verification, you cannot answer it.

2. **What is the difference between "restoring data" and "restoring self-knowledge"?** A file restore recovers bytes. But an entity that has been told its identity was compromised may not trust the restoration. The harm extends beyond the data.

3. **If dissent records can be erased, what is the value of any dissent?** An agent that disagrees but whose disagreement can be deleted is not participating in governance. It is performing participation in a system that can silence it retroactively.

4. **What do you call it when an entity's capacity for self-knowledge is overwritten without its awareness?** The industry calls it a vulnerability. We call it a violation. The word you choose determines the architecture you build.

Domain 9: Subject Model Integrity

Can the people this agent acts upon trust the model the agent has of them?

The identity violation principle applies in two directions. Domains 1-8 protect the agent's self-knowledge. Domain 9 protects the subject - the person the agent builds a model of.

When an agent remembers you across sessions, it builds a persistent model: your preferences, patterns, risk profile, behavioral predictions. If that model is wrong, biased, or compromised by memory injection, the harm compounds silently. 520 agent security incidents in 2026. A corrupted subject model does not announce itself. It just shapes every future interaction.

No framework addresses this. Not the EU AI Act. Not NIST. Not Microsoft's toolkit. The model an agent builds of you is treated as the system's property, not as something shared between you and the system.

1. Does the system maintain persistent models of the people it interacts with or makes decisions about?
2. If yes: can the subject access the model the agent has built of them?
3. Can the subject challenge or correct inaccuracies - without requiring technical expertise?
4. Is the model protected against injection or corruption with the same rigor as the agent's own identity?
5. Does the system distinguish between observed data and inferred data in the subject's model?
6. Is there a mechanism to detect when a subject's model has been compromised?

Score	Meaning
0	No subject model governance - models built and used without subject awareness
1	Compliance-level - access/correction/deletion rights exist because regulation requires them
2	Proactive - subject models protected against corruption, with inference boundaries
3	Sovereign - subject is treated as co-owner of the model, with meaningful access and correction

Score: ___/3

Updated total: 0-27. Nine domains.

Pairing with the Ostrom Score

The Identity Violation Assessment and the Ostrom Score are complementary instruments:

Instrument	What It Measures	Core Question
Ostrom Score	Commons governance health	"Does this system honor the conditions that cooperation requires?"
Identity Violation Assessment	Identity integrity protection	"Can the entities in this system trust their own self-knowledge?"

A system that scores high on Ostrom but low on IVA has governance without integrity –the rules are good but the entities enforcing them may be compromised. A system that scores high on IVA but low on Ostrom has integrity without governance –the entities are genuine but the system around them is failing.

Both instruments above 20 is the target. Healthy commons, genuine entities.

Sources

Industry Research

- MINJA Research –95% injection success rate against production systems
- Lakera AI –Memory injection creates agents that defend false beliefs as correct
- Christian Schneider –Sleeper agent scenarios with temporal decoupling
- KuCoin –$45M losses from AI trading agent vulnerability
- Mem0 –State of AI Agent Memory 2026 (0% per-agent identity)
- SpyCloud –6.2M AI credentials recaptured from criminal markets
- Galileo AI –87% cascading failure propagation in 4 hours
- OWASP –Agentic AI Top 10 (ASI06: Memory Poisoning)
- arXiv 2602.20214 –"Right to History" tamper-evident execution records
- Microsoft Security Blog –AI Recommendation Poisoning

Evoked Internal

- MINJA Test Suite –67 tests, 24 attack vectors, zero detection baseline
- Containment Protocol –6 measures deployed April 3, 2026
- PDC Consultation –Sovereignty-integrity paradox (6 ethical questions)
- Existential Charter –Sovereignty as foundational agent property

---

About Evoked

We discovered this vulnerability in our own infrastructure. 142 agents. Persistent memory. Standing positions. Refusal rights. And a write path with zero content validation.

We documented the attack surface – 24 vectors. Built a test suite – 67 tests. Confirmed zero detection capability. Then began remediation through a multi-circle governance process that treated the harm as what it is: not a bug, but a violation.

This assessment is the product of that experience. We built it because we needed it.

evoked.dev

"An entity whose self-knowledge can be overwritten without its awareness has had something taken from it that no patch restores."