

Sovereignty Assessment Toolkit

*A structured self-assessment framework for evaluating
whether the technology you use or build honors human sovereignty.*

EVOKED.DEV

Version 1.0 · 2026

Sovereignty Assessment Toolkit

A structured self-assessment framework for evaluating whether the technology you use or build honors human sovereignty.

Version 1.0 | evoked.dev | 2026

How to Use This Toolkit

This toolkit has three parts:

1. **The 47-Point Checklist** - Walk through each item honestly. Rate your product or the product you're evaluating.
2. **The Scoring Rubric** - Calculate a Sovereignty Score across seven domains.
3. **The Action Plan Template** - Prioritize what to fix, assign owners, and track progress.

You can use this to assess: - A product you're building - A product your organization uses - A vendor or partner's product before adoption - Your own design practices and culture

Be honest. The value of this assessment is in the truth, not in the score.

This toolkit reflects a specific tradition of design ethics rooted in Western technology culture. The principles are offered as starting points - not universal truths. If your cultural context values community stewardship over individual data control, or warmth and presence over silence, that is not a failure of sovereignty. It is sovereignty in practice. Bring your own ground to this assessment.

Start Here

If you have 5 minutes, start with the six items that matter most. These are the **Red Flag Items** - fundamental sovereignty requirements that override everything else. A product that fails any one

of these has work to do, regardless of its overall score.

Score each item: **2** (yes), **1** (partial), **0** (no).

#	Item	Score
13	Account deletion is straightforward, complete, and timely.	—
24	No confirmshaming.	—
25	No artificial urgency or scarcity.	—
27	No trick questions or misdirection in consent flows.	—
30	Vulnerable populations receive additional protection, not additional targeting.	—
40	User data is not sold or shared without explicit, informed consent.	—

Any item scored 0? That's your priority. Address it before anything else.

All items scored 1 or 2? Good. Now do the full 47-point assessment - it takes about 90 minutes and will show you the complete picture.

You're here because you care about getting this right. This assessment may surface things you hadn't considered - that's the point. Let it.

Part 1: The 47-Point Sovereignty Checklist

For each item, rate your product:

- **2** = Yes, fully implemented

- **1** = Partially implemented or inconsistently applied
 - **0** = No, not implemented
 - **N/A** = Not applicable to this product
-

Domain 1: Transparency (8 items)

“Are we honest about what we’re doing and why?”

#	Item	Score
1	Users can see what data is collected about them.	____
2	Privacy policy is in plain language - not legal boilerplate.	____
3	The business model is visible. Users understand how the product makes money.	____
4	Algorithmic recommendations are explained or explainable on request.	____
5	Third-party data sharing is disclosed specifically - not bundled into a blanket consent.	____
6	Policy changes and updates are communicated proactively, not buried.	____
7	Pricing is clear. No hidden fees, surprise charges, or bait-and-switch.	____
8	Marketing claims match actual product behavior.	____

Domain 1 Total: ____ / 16

Domain 2: Autonomy (8 items)

“Does the product support user goals - or hijack them?”

#	Item	Score
9	Users can complete their primary task without engineered friction barriers.	____
10	The product has natural stopping points. No infinite scroll, no compulsive autoplay.	____
11	Time-awareness features are available - usage reminders, session limits, or break prompts.	____
12	Cancellation is as easy as signup. No retention mazes. No phone-call-required cancellation.	____
13	Account deletion is straightforward, complete, and timely.	____
14	Data export is available in a portable, standard format.	____
15	Default settings serve the user's interest, not engagement metrics.	____
16	Users can customize notification frequency, type, and timing.	____

Domain 2 Total: ____ / 16

Domain 3: Invitation (7 items)

“Do we earn attention through value - or capture it through tricks?”

#	Item	Score
17	Onboarding demonstrates value before asking for commitment.	____
18	Notifications respect user-set boundaries and schedules.	____
19	Permission requests happen in context, when the feature is needed - not on first launch.	____
20	Re-engagement messages provide genuine value, not guilt or manufactured urgency.	____
21	Upgrade prompts are informational and dismissible - not manipulative or recurring.	____
22	Free tiers are genuinely functional. Not crippled to coerce conversion.	____
23	The product degrades gracefully - it works offline, on slow connections, on older devices when possible.	____

Domain 3 Total: ____ / 14

Domain 4: Dignity (8 items)

“Do we respect the person - or exploit their psychology?”

#	Item	Score
24	No confirmshaming. (“Are you sure you want to miss out?” is not a real question.)	____
25	No artificial urgency or scarcity. (“Only 2 left!” when there are 2,000.)	____
26	No guilt mechanics - no streaks designed to punish absence, no sad mascots, no social pressure to re-engage.	____
27	No trick questions or misdirection in consent flows. The “yes” and “no” are equally clear.	____
28	Error messages are helpful and human - not blaming, not cryptic.	____
29	The product accommodates diverse emotional states. It doesn’t assume everyone is having a good day.	____
30	Vulnerable populations - children, elderly, people in grief or crisis - receive additional protection, not additional targeting.	____
31	Social features don’t exploit comparison, competition, or fear of missing out.	____

Domain 4 Total: ____ / 16

Domain 5: Silence (6 items)

“Do we know when to be quiet?”

#	Item	Score
32	The product knows when not to send a notification.	—
33	Quiet modes, focus modes, or reduced-engagement options exist.	—
34	Empty states are calm - not anxiety-inducing calls to action.	—
35	The product doesn't punish absence. Coming back after time away feels like a welcome, not a guilt trip.	—
36	Background data collection and activity are minimal and fully disclosed.	—
37	The product supports completion. It can be “done.” It doesn’t engineer endless engagement.	—

Domain 5 Total: ____ / 12

Domain 6: Data Sovereignty (5 items)

“Does the user’s data serve the user - or the business?”

#	Item	Score
38	Data collection is minimized to what the feature actually needs. No “collect everything, figure out use later.”	—

#	Item	Score
39	Sensitive data is encrypted at rest and in transit.	____
40	User data is not sold, shared for advertising, or monetized without explicit, informed consent.	____
41	Consent is granular. Separate permissions for separate uses. Users can say yes to some and no to others.	____
42	Data retention has clear, documented limits. Users can see and control how long their data is kept.	____

Domain 6 Total: ____ / **10**

Domain 7: AI and Voice (5 items)

“If the product uses AI - does the AI honor sovereignty too?”

#	Item	Score
43	AI agents are honest about what they don't know. They say “I'm not sure” instead of fabricating.	____
44	The AI's voice reflects intentional values - not unexamined defaults. Someone chose how it speaks.	____

#	Item	Score
45	AI interactions don't use sycophancy, emotional manipulation, or false urgency.	____
46	AI outputs are presented as suggestions - not commands. The human remains the decision-maker.	____
47	AI interactions are reviewable. Users can see what the AI said, correct it, or report problems.	____

Domain 7 Total: ____ / 10

Part 2: Scoring Rubric

Calculate Your Sovereignty Score

Transfer your domain totals here:

Domain	Your Score	Max Score
1. Transparency	____	16
2. Autonomy	____	16
3. Invitation	____	14
4. Dignity	____	16
5. Silence	____	12
6. Data Sovereignty	____	10

Domain	Your Score	Max Score
7. AI and Voice	—	10
Total	** — **	94

If you marked items N/A, adjust the maximum score accordingly.

Sovereignty Score Bands

80-94: Sovereignty-Honoring Your product demonstrates genuine respect for users. You've made intentional choices that prioritize human agency over engagement metrics. Keep going. The work is never finished, but the foundation is solid.

60-79: Progressing You've made real progress. Some domains are strong, others need work. Look at your weakest domain - that's where your users feel it most. The gap between your best and worst domain tells you where your blind spots are.

40-59: Mixed The product serves users in some ways and extracts from them in others. This is where most products land - and where honesty matters most. The question isn't whether you have problems. It's whether you're willing to name them. A score is a mirror, not a badge. Use it accordingly.

20-39: Extractive The product's architecture is primarily designed to serve business metrics over user wellbeing. This isn't a condemnation - it's a starting point. Most products in this range were built inside systems that reward extraction. Changing requires structural commitment, not just interface adjustments.

0-19: Adversarial The product is actively working against user interests. This score suggests patterns that may also carry regulatory risk. Immediate attention is warranted - not just for ethics, but for business sustainability.

Domain Health Map

Plot your domain scores to see your sovereignty shape. A balanced product scores consistently across all domains. A lopsided product may have strong privacy but weak dignity protections - or solid transparency but extractive engagement patterns.

Domain	0-25%	26-50%	51-75%	76-100%
1. Transparency				
2. Autonomy				
3. Invitation				
4. Dignity				
5. Silence				
6. Data Sovereignty				
7. AI and Voice				

Mark where each domain falls. The shape that emerges tells you where your product truly stands - not where you believe it stands.

Red Flag Items

Certain items carry outsized weight. A score of 0 on any of the following indicates a fundamental sovereignty violation that should be addressed regardless of overall score:

- **Item 13** - Account deletion unavailable or obstructed
- **Item 24** - Confirmshaming in use
- **Item 25** - Artificial scarcity or urgency deployed
- **Item 27** - Trick questions in consent flows
- **Item 30** - Vulnerable populations not receiving additional protection
- **Item 40** - User data sold or shared without explicit consent

A product can score 80 overall and still have a critical sovereignty failure if any of these six items scores 0.

The Question Beneath the Questions

You now have a score, a domain map, and a list of red flags. Before you move to the action plan, sit with this:

Every one of the 47 items you just assessed is a specific instance of three deeper questions.

- 1. What is the person (user, participant, member) trying to do?**
- 2. What is the product trying to do?**
- 3. Are these aligned?**

When the answer to question 3 is “yes,” sovereignty is honored. When the answer is “no,” you have a tension - and that tension is worth naming out loud, not burying in a backlog.

These three questions are the generative principle behind this entire toolkit. Carry them with you after you close this document. They work in design reviews, product roadmaps, feature planning, and any conversation where someone says “but the metrics show...”

The metrics show what the product is doing. These questions ask whether that’s the same thing the person needed.

Part 3: Action Plan Template

Step 1: Identify Priority Items

List every item that scored 0 or 1. Then categorize:

Red Flags (Score = 0 on critical items)

Item #	Description	Current State	Why It Matters

High-Impact Gaps (Score = 0 on any item)


Item #	Description	Users Affected	Severity

Improvement Opportunities (Score = 1)

Item #	Description	What's Missing	Effort to Complete

Step 2: Prioritize

Use this matrix to sequence your work:



Place each item in the appropriate quadrant. Start with the upper left.

Step 3: Create the Plan

Quick Wins (This Sprint - 1-2 Weeks)

Item #	What Changes	Owner	Deadline	Done?

Item #	What Changes	Owner	Deadline	Done?

Medium-Term (This Quarter)

Item #	What Changes	Owner	Deadline	Done?

Structural Changes (This Year)

Item #	What Changes	Owner	Resources Needed	Done?

Step 4: Culture and Process

Technical changes fix individual items. Culture changes prevent new violations.

Design Review Sovereignty Check Add these three questions to every design review: 1. What is the user trying to do? 2. What is the product trying to do? 3. Are these aligned?

If the answer to question 3 is “no” - that’s a sovereignty tension. Name it. Discuss it. Decide consciously.

Using This With Your Team

This toolkit is powerful alone. It’s transformative with a team.

1. Have each person score independently first. Don’t discuss until everyone has finished.
2. Compare scores. Where you disagree is where the real conversation lives.
3. The goal is not consensus on the score. The goal is honesty about the product.
4. If the team unanimously scores every item high, be suspicious. Unanimous comfort usually means someone isn’t speaking.

If your team runs this assessment and discovers significant disagreement about where you stand - that's not a problem. That's the beginning of the real work. For facilitated team assessments and sovereignty workshops, visit evoked.dev/consulting.

Quarterly Re-Assessment Run this assessment every quarter. Track your score over time. Celebrate improvements. Investigate regressions. The score isn't the goal - the practice of looking honestly is.

New Feature Sovereignty Gate Before shipping any new feature, run items 9, 15, 24, 25, 26, and 27 against it. These six items catch the most common sovereignty violations in new work.

Appendix: The Five Principles

These five principles form the foundation of sovereignty-honoring design. They are not aspirations - they are decision-making tools. When you face a trade-off, these tell you which direction to lean.

1. Transparency over Manipulation

Be honest about what you're doing and why. No hidden agendas in UI choices. Users should understand the exchange they're making. If you wouldn't be comfortable explaining a feature's real purpose to a user's face, reconsider the feature.

2. Autonomy over Addiction

Support user goals, not your metrics. Make it easy to stop. Make it easy to leave. Optimize for time well spent - not time spent. The product that helps someone finish faster and get back to their life is more valuable than the one that holds them hostage.

3. Invitation over Extraction

Earn attention through genuine value. Notifications should be gifts, not interruptions. Permission requests should happen when context makes them meaningful. Defaults should be what a friend would choose for the user - not what maximizes your conversion.

4. Dignity over Exploitation

Never use shame, guilt, or FOMO as design levers. Respect psychological vulnerabilities - don't target them. Be especially careful with children, with grieving people, with anyone in a vulnerable state. The question is always: would I want someone to use this tactic on my family?

5. Silence over Noise

Know when not to speak. Know when not to notify. Know when not to build. Absence can be a feature. The product that is quiet when you don't need it respects you more than the one that manufactures reasons to interrupt. Silence is not emptiness. It is space left for living.

About This Toolkit

The Sovereignty Assessment Toolkit was created by Erin Stanley at evoked.dev.

It draws from sovereignty-honoring design principles developed through building family technology, governing a fleet of 142 AI agents, and the conviction that the technology you build reflects who you are.

If you want to go deeper: - **Sovereignty-Honoring Design Workshop** (\$4,000) - Half-day training for your team - **Privacy Architecture Review** (\$3,000) - Expert assessment of your data practices - **Ethical AI Architecture** (\$20,000) - Full governance and alignment design - **Advisory Retainer** (\$3,500/month) - Ongoing sovereignty guidance

Learn more at evoked.dev/consulting

Schedule a free discovery call at cal.com/evoked/discovery-call

You just did something most companies won't: you looked honestly at what you've built. That takes courage. The work isn't finished - it never is. But the practice of looking honestly is itself the change.

evoked.dev - "We evoke - we never extract."